

project-pdf

November 1, 2023

1 Complete the steps given below:

- 1.1 1. Perform necessary EDA on the data
- 1.2 2. Use machine learning to create clusters of similar projects.
- 1.3 3. Create a regression model to predict the budget
- 1.4 4. Create a classification model to predict the value of the Type column.

```
[1]: import pandas as pd

df = pd.read_csv('Freelance_Platform_Projects[1].csv') # reading csv file
df.head() # display first five rows
```

```
[1]:
```

	Title	Category Name	Experience \
0	Banner images for web desgin websites	Design	Entry (\$)
1	Make my picture a solid silhouette	Video, Photo & Image	Entry (\$)
2	Bookkeeper needed	Business	Entry (\$)
3	Accountant needed	Business	Entry (\$)
4	Guest Post on High DA Website	Digital Marketing	Expert (\$\$\$)

	Sub Category Name	Currency	Budget	Location \
0	Graphic Design	EUR	60.0	remote
1	Image Editing	GBP	20.0	remote
2	Finance & Accounting	GBP	12.0	remote
3	Tax Consulting & Advising	GBP	14.0	remote
4	SEO	USD	10000.0	remote

	Freelancer Preferred From	Type	Date Posted \
0	ALL	fixed_price	2023-04-29 18:06:39
1	ALL	fixed_price	2023-04-29 17:40:28
2	ALL	fixed_price	2023-04-29 17:40:06
3	ALL	fixed_price	2023-04-29 17:32:01
4	ALL	fixed_price	2023-04-29 17:09:36

	Description	Duration \
0	We are looking to improve the banner images on...	NaN
1	Hello \n\nI need a quick designer to make 4 pi...	NaN

```

2 Hi - I need a bookkeeper to assist with bookke...      NaN
3 Hi - I need an accountant to assist me with un...      NaN
4 Hi, I am currently running a project where I w...      NaN

```

```

      Client Registration Date Client City Client Country Client Currency \
0          2010-11-03      Dublin      Ireland      EUR
1          2017-02-21      London  United Kingdom      GBP
2          2023-04-09      London  United Kingdom      GBP
3          2023-04-09      London  United Kingdom      GBP
4          2016-07-01      Mumbai      India      USD

```

```

      Client Job Title
0      PPC Management
1      Office manager
2      Paralegal
3      Paralegal
4 Guest posts buyer

```

```
[2]: df.tail() # display last five rows
```

```

[2]:
      Title \
12217 Published Travel Writer required for content c...
12218 Shopify - Filtering Work (Product Selection/No...
12219      Simple SQL Query
12220 Create a Carbon, Water, Waste Calculating plat...
12221      COMPANY REGISTERS

```

```

      Category Name      Experience      Sub Category Name \
12217 Writing & Translation      Entry ($)      Content Writing
12218      Design  Intermediate ($$)      Web Design
12219 Technology & Programming      Entry ($)  Data Science & Analysis
12220      Design      Expert ($$$)      Web Design
12221      Business      Expert ($$$)  Administration Assistance

```

```

      Currency Budget      Location Freelancer Preferred From      Type \
12217      GBP   50.0      remote      ALL fixed_price
12218      GBP   65.0  remote_country      GB fixed_price
12219      GBP   50.0      remote      ALL fixed_price
12220      USD   39.0      remote      ALL      hourly
12221      GBP   75.0      remote      ALL fixed_price

```

```

      Date Posted      Description \
12217 2023-01-18 19:23:01 I am looking for a published travel writer to ...
12218 2023-01-18 19:18:48 On our website www.juicebitz.co.uk we have add...
12219 2023-01-18 19:18:48 I need someone to write a quick SQL query on a...
12220 2023-01-18 19:18:47 I am seeking a full stack web developer who sp...
12221 2023-01-18 19:18:47 Hi, the following administrative task would be...

```

	Duration	Client Registration Date	Client City	Client Country	\
12217	NaN	2011-06-06	Amsterdam	Netherlands	
12218	1 day or less	2022-03-23	Filey	United Kingdom	
12219	NaN	2022-03-14	London	United Kingdom	
12220	NaN	2013-07-21	Noida	India	
12221	NaN	2020-09-21	Grays	United Kingdom	

	Client Currency	Client Job Title
12217	GBP	Wordpress Expert
12218	GBP	Director
12219	GBP	NaN
12220	USD	Google Adwords, Pay Per Click, Google Shopping...
12221	GBP	NaN

```
[3]: df.size # no of total elements
```

```
[3]: 207774
```

```
[4]: df.shape # no of rows = 12222 , no of column = 17
```

```
[4]: (12222, 17)
```

```
[5]: df.info() # detailed info of your data
```

```
# Number of rows = 12222
# Number of columns = 17
# For every column
# => Name of column
# => Number of not null values
# => Number of null value = Total rows - Not null values
# => Data type
# Number of columns for each data type
# Memory usage
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 12222 entries, 0 to 12221
```

```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	Title	12222 non-null	object
1	Category Name	12222 non-null	object
2	Experience	12222 non-null	object
3	Sub Category Name	12222 non-null	object
4	Currency	12222 non-null	object
5	Budget	12222 non-null	float64
6	Location	12222 non-null	object

```

7  Freelancer Preferred From  12222 non-null  object
8  Type                      12222 non-null  object
9  Date Posted              12222 non-null  object
10 Description              12222 non-null  object
11 Duration                 1602 non-null   object
12 Client Registration Date  12222 non-null  object
13 Client City              12222 non-null  object
14 Client Country           12222 non-null  object
15 Client Currency          12222 non-null  object
16 Client Job Title         4588 non-null   object
dtypes: float64(1), object(16)
memory usage: 1.6+ MB

```

```
[6]: df.drop(['Date Posted', 'Client Registration Date'], axis = 1, inplace = True)
      ↪# dropping unnecessary columns
```

2 1) EDA

Perform necessary EDA on the data

2.1 Handling Missing Values

```
[7]: df.isnull().sum() # total no of missing values in each column
```

```

[7]: Title                      0
     Category Name              0
     Experience                 0
     Sub Category Name          0
     Currency                   0
     Budget                     0
     Location                   0
     Freelancer Preferred From  0
     Type                       0
     Description                 0
     Duration                   10620
     Client City                0
     Client Country             0
     Client Currency            0
     Client Job Title           7634
     dtype: int64

```

Here, we have 2 option to deal with this: 1) Deletion : removal of data or records from a dataset.
2) Imputation : process of filling in missing values with estimated or substituted data.

```

[8]: # display the rows in which the 'Duration' column has missing values.

missing_rows = df.index[df['Duration'].isna()==True]

```

```
df.loc[missing_rows, 'Duration']
```

```
[8]: 0      NaN
      1      NaN
      2      NaN
      3      NaN
      4      NaN
      ...
     12216   NaN
     12217   NaN
     12219   NaN
     12220   NaN
     12221   NaN
      Name: Duration, Length: 10620, dtype: object
```

```
[9]: # Imputation (Maximum values are missing so we can't delete it because after
      ↪ deletion we will have very less data)

      df['Duration'].fillna("Not Available", inplace=True)
      df.loc[missing_rows,:]
```

```
[9]:
```

	Title \
0	Banner images for web desgin websites
1	Make my picture a solid silhouette
2	Bookkeeper needed
3	Accountant needed
4	Guest Post on High DA Website
...	...
12216	Simple SQL Query
12217	Published Travel Writer required for content c...
12219	Simple SQL Query
12220	Create a Carbon, Water, Waste Calculating plat...
12221	COMPANY REGISTERS

	Category Name	Experience	Sub Category Name \
0	Design	Entry (\$)	Graphic Design
1	Video, Photo & Image	Entry (\$)	Image Editing
2	Business	Entry (\$)	Finance & Accounting
3	Business	Entry (\$)	Tax Consulting & Advising
4	Digital Marketing	Expert (\$\$\$)	SEO
...
12216	Technology & Programming	Entry (\$)	Databases
12217	Writing & Translation	Entry (\$)	Content Writing
12219	Technology & Programming	Entry (\$)	Data Science & Analysis
12220	Design	Expert (\$\$\$)	Web Design
12221	Business	Expert (\$\$\$)	Administration Assistance

	Currency	Budget	Location	Freelancer Preferred From	Type \
0	EUR	60.0	remote	ALL	fixed_price
1	GBP	20.0	remote	ALL	fixed_price
2	GBP	12.0	remote	ALL	fixed_price
3	GBP	14.0	remote	ALL	fixed_price
4	USD	10000.0	remote	ALL	fixed_price
...
12216	GBP	30.0	remote	ALL	fixed_price
12217	GBP	50.0	remote	ALL	fixed_price
12219	GBP	50.0	remote	ALL	fixed_price
12220	USD	39.0	remote	ALL	hourly
12221	GBP	75.0	remote	ALL	fixed_price

	Description	Duration \
0	We are looking to improve the banner images on...	Not Available
1	Hello \n\nI need a quick designer to make 4 pi...	Not Available
2	Hi - I need a bookkeeper to assist with bookke...	Not Available
3	Hi - I need an accountant to assist me with un...	Not Available
4	Hi, I am currently running a project where I w...	Not Available
...
12216	I need someone to write a quick SQL query base...	Not Available
12217	I am looking for a published travel writer to ...	Not Available
12219	I need someone to write a quick SQL query on a...	Not Available
12220	I am seeking a full stack web developer who sp...	Not Available
12221	Hi, the following administrative task would be...	Not Available

	Client City	Client Country	Client Currency \
0	Dublin	Ireland	EUR
1	London	United Kingdom	GBP
2	London	United Kingdom	GBP
3	London	United Kingdom	GBP
4	Mumbai	India	USD
...
12216	London	United Kingdom	GBP
12217	Amsterdam	Netherlands	GBP
12219	London	United Kingdom	GBP
12220	Noida	India	USD
12221	Grays	United Kingdom	GBP

	Client Job Title
0	PPC Management
1	Office manager
2	Paralegal
3	Paralegal
4	Guest posts buyer
...	...
12216	NaN

```

12217                                Wordpress Expert
12219                                NaN
12220  Google Adwords, Pay Per Click, Google Shopping...
12221                                NaN

```

```
[10620 rows x 15 columns]
```

```
[10]: # display the rows in which the 'Client Job Title' column has missing values.
```

```

missing_rows = df.index[df['Client Job Title'].isna()==True]
df.loc[missing_rows, 'Client Job Title']

```

```

[10]: 6      NaN
      8      NaN
      9      NaN
      11     NaN
      15     NaN

```

```

...
12213     NaN
12214     NaN
12216     NaN
12219     NaN
12221     NaN

```

```
Name: Client Job Title, Length: 7634, dtype: object
```

```
[11]: # Imputation
```

```

df['Client Job Title'].fillna("Not Available", inplace=True)
df.loc[missing_rows, :]

```

```

[11]:
      6      Make web site for Tutoring company
      8      E-learning
      9      19 sentences recording. native english speaker...
     11      Looking for someone to configure AWS server
     15      Logo colour change
...
12213  Need someone with Clickhouse Database expertise
12214      modify existing flyer
12216      Simple SQL Query
12219      Simple SQL Query
12221  COMPANY REGISTERS

```

	Category Name	Experience	Sub Category Name \
6	Design	Entry (\$)	Web Design
8	Design	Intermediate (\$\$)	Illustration & Drawing
9	Music & Audio	Entry (\$)	Voice-Over

11	Technology & Programming	Expert (\$\$\$)	Website Development
15	Design	Entry (\$)	Logo Design
...
12213	Technology & Programming	Entry (\$)	Databases
12214	Design	Entry (\$)	Graphic Design
12216	Technology & Programming	Entry (\$)	Databases
12219	Technology & Programming	Entry (\$)	Data Science & Analysis
12221	Business	Expert (\$\$\$)	Administration Assistance

	Currency	Budget	Location	Freelancer Preferred From	Type \
6	USD	10.0	remote	ALL	fixed_price
8	GBP	0.0	remote	ALL	fixed_price
9	USD	10.0	remote	ALL	fixed_price
11	EUR	31.0	remote	ALL	hourly
15	USD	10.0	remote	ALL	fixed_price
...
12213	USD	120.0	remote	ALL	fixed_price
12214	USD	15.0	remote	ALL	fixed_price
12216	GBP	30.0	remote	ALL	fixed_price
12219	GBP	50.0	remote	ALL	fixed_price
12221	GBP	75.0	remote	ALL	fixed_price

	Description	Duration \
6	I need to build web site for my tutoring compa...	1 - 2 weeks
8	Looking for a quote for an introductory e-lear...	Not Available
9	1. need native speaker from US or UK or CA\n2...	Not Available
11	Hi we are looking to deploy our domain from go...	Not Available
15	Very easy job, I literally just need my logo c...	Not Available
...
12213	I want to do a clickhouse database project whi...	Not Available
12214	i would like to modify a existing flyer which ...	Not Available
12216	I need someone to write a quick SQL query base...	Not Available
12219	I need someone to write a quick SQL query on a...	Not Available
12221	Hi, the following administrative task would be...	Not Available

	Client City	Client Country	Client Currency	Client Job Title
6	London	United Kingdom	USD	Not Available
8	Glasgow	United Kingdom	GBP	Not Available
9	Sydney	Australia	GBP	Not Available
11	Budapest	Hungary	USD	Not Available
15	Menai	Australia	USD	Not Available
...
12213	Fremont	United States	USD	Not Available
12214	Dieburg	Germany	EUR	Not Available
12216	London	United Kingdom	GBP	Not Available
12219	London	United Kingdom	GBP	Not Available
12221	Grays	United Kingdom	GBP	Not Available


[7634 rows x 15 columns]

2.2 Handling Outliers

```
[12]: import seaborn as sns


sns.boxplot(df['Budget']) # checking outliers (we are using 'Budget' column
↪ because it is only numeric column we have)
```

[12]: <Axes: >



```
[13]: import matplotlib.pyplot as plt

plt.hist(df['Budget']) # Create a histogram of the 'Budget' column in the
↪ DataFrame 'df'
plt.show()
```



```
[14]: # probably the reason of so many outliers is 'skewness'
```

```
df['Budget'].skew() # >5 is considered as high
```

```
[14]: 42.455398395555996
```

```
[15]: # to handle skewness we use boxcox
```

```
from scipy.stats import boxcox
```


```
df['Budget'] = df['Budget'] + 1 # Adding 1 to make values positive
```

```
df['Budget'] = boxcox(df['Budget'])[0]
```

```
df['Budget'].skew()
```

```
[15]: -0.021454219174626175
```

```
[16]: plt.hist(df['Budget'])  
plt.show()
```




```
[17]: # Calculate the IQR, lower, and upper bounds

Q1 = df['Budget'].quantile(0.25)
Q3 = df['Budget'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Identify and remove outliers
outliers = df[(df['Budget'] < lower_bound) | (df['Budget'] > upper_bound)]
df = df[(df['Budget'] >= lower_bound) & (df['Budget'] <= upper_bound)]

# Create the box plot
sns.boxplot(df['Budget'])
```

```
[17]: <Axes: >
```



2.3 Categorical Data Encoding

```
[18]: df.dtypes
```

```
[18]: Title                object
      Category Name        object
      Experience            object
      Sub Category Name     object
      Currency              object
      Budget                float64
      Location              object
      Freelancer Preferred From object
      Type                  object
      Description            object
      Duration              object
      Client City           object
      Client Country        object
      Client Currency       object
      Client Job Title      object
      dtype: object
```

```
[19]: # Label Encoding (Label encoding is used to represent categorical data as
      ↪ ordinal integers)

from sklearn.preprocessing import LabelEncoder

def label_encode_columns(df, columns):
    for col in columns:
        encoder = LabelEncoder()
        encoding = encoder.fit_transform(df[col])
        df[col] = encoding
        var = df[col].head()
        print(var)

# List of columns to label encode
columns_to_encode = ['Title', 'Category Name', 'Experience', 'Sub Category',
                    ↪ 'Name', 'Currency',
                        'Location', 'Freelancer Preferred From', 'Type',
                    ↪ 'Description', 'Duration',
                        'Client City', 'Client Country', 'Client Currency',
                    ↪ 'Client Job Title']

# Call the function to label encode the specified columns
label_encode_columns(df, columns_to_encode)
```

```
0      956
1     6335
2     1094
3       461
5     1800
Name: Title, dtype: int32
0       1
1       7
2       0
3       0
5       6
Name: Category Name, dtype: int32
0       0
1       0
2       0
3       0
5       1
Name: Experience, dtype: int32
0      42
1      45
2      37
3      90
5      26
```

```

Name: Sub Category Name, dtype: int32
0    0
1    1
2    1
3    1
5    0
Name: Currency, dtype: int32
0    1
1    1
2    1
3    1
5    1
Name: Location, dtype: int32
0    1
1    1
2    1
3    1
5    1
Name: Freelancer Preferred From, dtype: int32
0    0
1    0
2    0
3    0
5    0
Name: Type, dtype: int32
0   10364
1    1236
2    2161
3    2163
5     564
Name: Description, dtype: int32
0    21
1    21
2    21
3    21
5    21
Name: Duration, dtype: int32
0    489
1    936
2    936
3    936
5    488
Name: Client City, dtype: int32
0     61
1    129
2    129
3    129
5    128

```

Name: Client Country, dtype: int32

```
0    0
1    1
2    1
3    1
5    0
```

Name: Client Currency, dtype: int32

```
0    1177
1    1127
2    1194
3    1194
5     149
```

Name: Client Job Title, dtype: int32

Scaling : Scaling improve the performance and interpretability of various machine learning algorithms. Types - 1) StandardScaler 2) MinMaxScaler

```
[20]: # we are using SS because it is suitable for a wide range of data and models

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df['Budget'] = scaler.fit_transform(df[['Budget']]) # fit - calculating mean
↳ and sd of data & transform - value after scaling
df['Budget'][:5]
```

```
[20]: 0    -0.123392
1    -1.141924
2    -1.639954
3    -1.488617
5     1.571911
Name: Budget, dtype: float64
```

```
[21]: df['Budget'].mean() # Calculate the mean value of the 'result' data
```

```
[21]: 8.809324647310018e-16
```

```
[22]: import numpy as np
median_value = np.median(df['Budget']) # Calculate the median value of the
↳ 'result' data using numpy
median_value
```

```
[22]: 0.12815933375328264
```

```
[23]: from scipy import stats
mode_value = stats.mode(df['Budget']) # Calculate the mode value of the
↳ 'result' data using scipy's mode function
mode_value
```

```
C:\Users\Sapna\AppData\Local\Temp\ipykernel_10196\2282642316.py:2:
```

```
FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy
1.11.0, this behavior will change: the default value of `keepdims` will become
False, the `axis` over which the statistic is taken will be eliminated, and the
value None will no longer be accepted. Set `keepdims` to True or False to avoid
this warning.
```

```
mode_value = stats.mode(df['Budget']) # Calculate the mode value of the
'result' data using scipy's mode function
```

```
[23]: ModeResult(mode=array([0.31861116]), count=array([743]))
```

```
[24]: df['Budget'].std() # Calculate the standard deviation of the 'result' data
```

```
[24]: 1.0000412056781718
```

```
[25]: df['Budget'].min() # Find the minimum value in the 'Budget' column of the
↳ DataFrame 'df'
```

```
[25]: -2.322457346222455
```

```
[26]: df['Budget'].max() # Find the maximum value in the 'Budget' column of the
↳ DataFrame 'df'
```

```
[26]: 2.7526630506423064
```

Unsupervised Learning : a type of machine learning where the algorithm is trained on data without explicit supervision or labeled outcomes.(most common way of doing UL is 'Clustering')

3 2) Cluster

Use machine learning to create clusters of similar projects. Data points within one cluster are expected to be similar to each other.

Data points in different clusters are expected to be different from each other.

k-mean clustering : It aims to partition the n observation into k ($\leq n$) sets so as to minimize the within cluster sum of squares(wcss)

```
[27]: x = df.iloc[:, [0,1,2,3,4]]
x.head()
```

```
[27]:
```

	Title	Category Name	Experience	Sub Category Name	Currency
0	956	1	0	42	0
1	6335	7	0	45	1
2	1094	0	0	37	1
3	461	0	0	90	1
5	1800	6	1	26	0


```
[28]: from sklearn.cluster import KMeans

wcscs_list = []
# Loop through a range of cluster numbers from 1 to 10
for i in range(1,11):
    model = KMeans(n_clusters=i)
    model.fit(x)
    wcscs = model.inertia_ # Calculate the WCSS for the current number of
    ↪clusters
    wcscs_list.append(wcscs) # Append the WCSS value to the list

wcscs_list # Display the list of WCSS values for different cluster numbers
```


```
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
```

```
1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
```

```
[28]: [133590593098.89883,
      33673762370.248634,
      14790098656.414959,
      8401938819.139595,
      5352760997.224634,
      3756407150.835943,
      2727560227.2335005,
      2106972434.0767555,
      1655536457.5855198,
      1357347335.5917444]
```

```
[29]: import matplotlib.pyplot as plt

plt.plot(range(1,11), wcss_list, '-o') # Create a line plot of the WCSS values,
    ↪ for different numbers of clusters
plt.title('Elbow Method Graph') # Set the title of the plot
plt.xlabel('Number of Clusters') # Set the label for the x-axis
plt.ylabel('WCSS') # Set the label for the y-axis
plt.show() # Display the plot
```



```
[30]: kmeans = KMeans(n_clusters=3)
      kmeans.fit(x) # no 'y' because it's unsupervised
      pred = kmeans.predict(x)
      pred[:3]
```

C:\Users\Sapna\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1412:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)

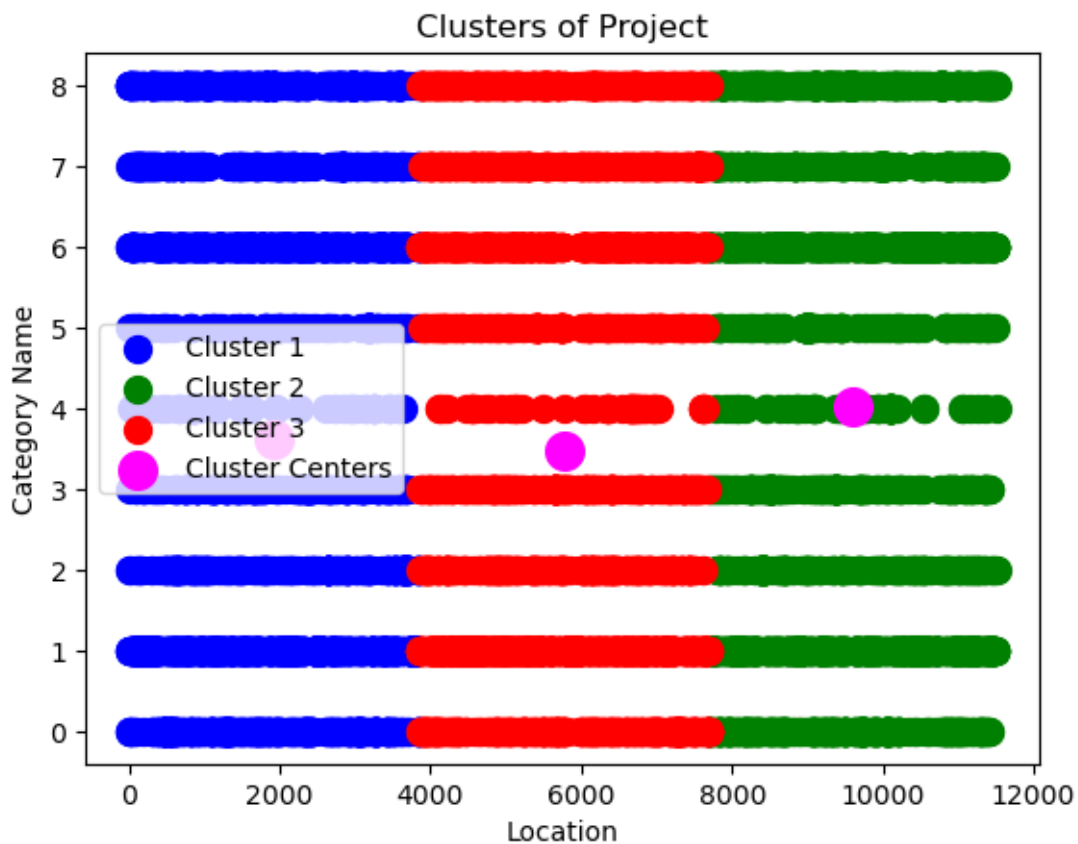
```
[30]: array([0, 2, 0])
```

```
[31]: kmeans.cluster_centers_
```

```
[31]: array([[1.92603392e+03, 3.61252785e+00, 7.01658826e-01, 5.25768755e+01,
            1.17504333e+00],
            [9.60057880e+03, 4.03201787e+00, 6.54753040e-01, 6.25624224e+01,
            1.17175478e+00],
            [5.78141775e+03, 3.47848537e+00, 6.95106958e-01, 5.51703959e+01,
            1.20752397e+00]])
```

```
[32]: # Create scatter plots for each cluster and cluster centers

plt.scatter(x.iloc[pred==0, 0], x.iloc[pred==0, 1],
            s=100, c='blue', label='Cluster 1')
plt.scatter(x.iloc[pred==1, 0], x.iloc[pred==1, 1],
            s=100, c='green', label='Cluster 2')
plt.scatter(x.iloc[pred==2, 0], x.iloc[pred==2, 1],
            s=100, c='red', label='Cluster 3')
plt.scatter(kmeans.cluster_centers_[0,0],
            kmeans.cluster_centers_[0, 1],
            s=200, c='magenta', label='Cluster Centers')
plt.title('Clusters of Project') # Set the title of the plot
plt.xlabel('Location') # Set the label for the x-axis
plt.ylabel('Category Name') # Set the label for the y-axis
plt.legend() # Add a legend to the plot to distinguish data points and cluster_
            ↪centers
plt.show() # Display the plot
```



4 3) Regression

Create a regression model to predict the 'budget'. Target variable is continuous('Budget')

```
[33]: df.head()
```

```
[33]:
```

	Title	Category	Name	Experience	Sub	Category	Name	Currency	Budget	\
0	956		1	0		42	0	-0.123392		
1	6335		7	0		45	1	-1.141924		
2	1094		0	0		37	1	-1.639954		
3	461		0	0		90	1	-1.488617		
5	1800		6	1		26	0	1.571911		

	Location	Freelancer	Preferred	From	Type	Description	Duration	\
0	1			1	0	10364	21	
1	1			1	0	1236	21	
2	1			1	0	2161	21	
3	1			1	0	2163	21	
5	1			1	0	564	21	

	Client	City	Client	Country	Client	Currency	Client	Job	Title
0		489		61		0			1177
1		936		129		1			1127
2		936		129		1			1194
3		936		129		1			1194
5		488		128		0			149

```
[34]: # split the data into x and y

x = df.drop(columns=['Budget'])
y = df['Budget']
```

```
[35]: # Split the data into training and testing

from sklearn.model_selection import train_test_split

xtrain, xtest, ytrain, ytest=train_test_split(x,y,
                                              train_size=0.7,
                                              random_state=1)
```

4.0.1 XGBoost Regression

```
[36]: !pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\sapna\anaconda3\lib\site-
packages (2.0.0)
Requirement already satisfied: numpy in c:\users\sapna\anaconda3\lib\site-
packages (from xgboost) (1.24.3)
```

Requirement already satisfied: scipy in c:\users\sapna\anaconda3\lib\site-packages (from xgboost) (1.10.1)

```
[37]: import xgboost as xgb
```

```
model = xgb.XGBRegressor() # Create an instance of the XGBRegressor model
model.fit(xtrain, ytrain) # Fit the model to the training data
trainpred = model.predict(xtrain) # Make predictions on the training data
trainpred[:5] # Display the first 5 predictions
```

```
[37]: array([ 1.3640537 ,  1.3072898 , -1.5222318 ,  0.99022496, -0.66387075],
      dtype=float32)
```

```
[38]: testpred = model.predict(xtest) # Make predictions on the test data using the
      ↪trained model
      testpred # Display the test predictions
```

```
[38]: array([-0.7291285 ,  0.70003724, -1.0932764 , ..., -0.944017 ,
      -0.6009989 , -1.0011668 ], dtype=float32)
```

```
[39]: from sklearn.metrics import mean_squared_error
```

```
mean_squared_error(ytrain, trainpred) # Calculate the mean squared error for
      ↪training data
```

```
[39]: 0.13185775409381667
```

```
[40]: mean_squared_error(ytest, testpred) # Calculate the mean squared error for test
      ↪data
```

```
[40]: 0.3316048886167132
```

```
[41]: from sklearn.metrics import mean_absolute_error
```

```
# Calculate the mean absolute error (MAE) for the training predictions and true
      ↪training labels
mae_train = mean_absolute_error(ytrain, trainpred)
mae_train
```

```
[41]: 0.25981870039243443
```

```
[42]: # Calculate the mean absolute error (MAE) for the testing predictions and true
      ↪testing labels
mae_test = mean_absolute_error(ytest, testpred)
mae_test
```

```
[42]: 0.42899353652810474
```

5 4) Classification

Create a classification model to predict the value of the 'Type' column.

```
[43]: # split the data into x and y
```

```
x = df.drop(columns=['Type'])
y = df['Type']
```

```
[44]: from sklearn.model_selection import train_test_split
```

```
xtrain, xtest, ytrain, ytest = train_test_split(x, y,
                                                train_size=0.7,
                                                random_state=1)
```

5.0.1 Random Forest Classification

```
[45]: from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier(n_estimators=100) # Create a Random Forest
↳ Classifier model with 100 trees (n_estimators)
model.fit(xtrain, ytrain) # Fit the model to the training data
trainpred = model.predict(xtrain) # Make predictions on the training data
trainpred[:5] # Display the first 5 predictions
```

```
[45]: array([0, 0, 1, 0, 1])
```

```
[46]: model.predict_proba(xtrain)[:5] # Predict probabilities for training data
```

```
[46]: array([[1.  , 0.  ],
            [1.  , 0.  ],
            [0.1 , 0.9 ],
            [1.  , 0.  ],
            [0.08, 0.92]])
```

```
[47]: from sklearn.metrics import classification_report
```

```
# Generate a classification report to evaluate the model's performance on the
↳ training data
print(classification_report(ytrain, trainpred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7241
1	1.00	1.00	1.00	1253
accuracy			1.00	8494
macro avg	1.00	1.00	1.00	8494

weighted avg	1.00	1.00	1.00	8494
--------------	------	------	------	------

```
[48]: # Generate a classification report to evaluate the model's performance on the
      ↪ test data
      testpred = model.predict(xtest)
      print(classification_report(ytest, testpred))
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	3109
1	0.88	0.75	0.81	532
accuracy			0.95	3641
macro avg	0.92	0.87	0.89	3641
weighted avg	0.95	0.95	0.95	3641